

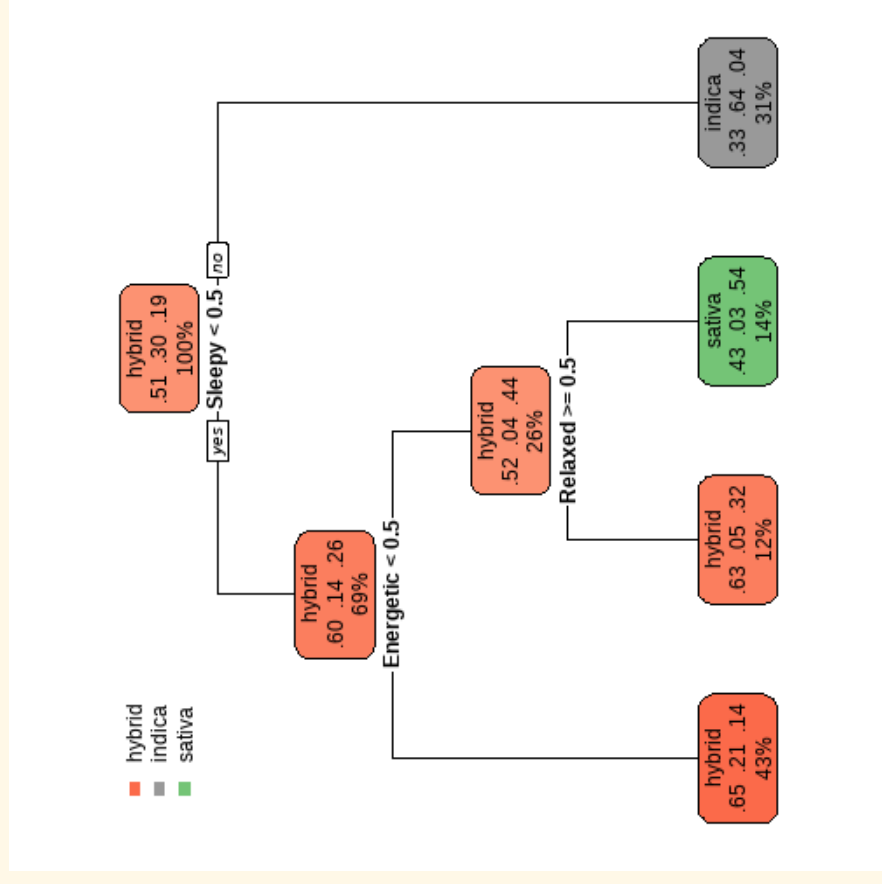
Decision Trees: Next Level

Measuring Success

Gini Index and Cost Complexity

Recall: Our simplest cannabis tree

- Which of the **final nodes** (or **leaves**) is **most pure**?
- Which is **least pure**?
-
- Could we split a node further for better purity?
- **almost certainly, yes!** It's highly unlikely that *all* of the unused variable have *exactly* the same prevalence across categories.
-
- Should we do it, or is that overfitting?



Cost Complexity revisited

- What is the **classification error** of each leaf?

```
rpart.plot(tree_fitted$fit)
```

- (Left to right)

- 0.35
- 0.37
- 0.46
- 0.36

- The **Gini Index** for a particular node is the average of errors in each class:

$$(0.35 * 0.65) + (0.21 * 0.79) + (0.14 * 0.86) = 0.5138$$

- + small values if the classification errors are close to 0, i.e., high node purity
- + large values (near 1) if the errors are high
- + this is related to the *variance* of the node

Gini Index

To calculate the Gini Index average across all leaves:

```
cann %>%  
  bind_cols(  
    predict(tree_fitted, cann, type = "prob")  
  ) %>%  
  gain_capture(truth = Type,  
               .pred_hybrid, .pred_indica, .pred_sativa)
```

```
## # A tibble: 1 x 3  
##   .metric      .estimator .estimate  
##   <chr>      <chr>      <dbl>  
## 1 gain_capture macro      0.482
```

Cost Complexity revisited

So, when should we split the tree further?

Only if the new splits improve the Gini Index by a certain amount.

This is the `cost_complexity` parameter!

--

But wait! This is a **penalized** metric, using an arbitrary penalty α to avoid overfitting.

Don't we like cross-validation better?

Well... yes.

But imagine fitting *every possible tree* and cross-validating.... yikes.

We have to limit our options and cut our losses somehow!

Bagging

Tree variability

Suppose I took two random subsamples of my cannabis dataset:

```
set.seed(9374534)
splits <- cann %>%
  initial_split(0.5, strata = Type)

cann_1 <- splits %>% training()
cann_2 <- splits %>% testing()
```

```
dim(cann_1)
```

```
## [1] 1154 69
```

```
dim(cann_2)
```

```
## [1] 1151 69
```

Tree variability

Then I fit a **decision tree** to each:

```
tree_1 <- tree_wflow %>%  
  fit(cann_1)  
  
tree_2 <- tree_wflow %>%  
  fit(cann_2)
```

How similar will the results be?

Tree variability

```
tree_1 %>%  
  pull_workflow_fit() %$%  
  fit %>%  
  rpart.plot()
```

Tree variability

```
tree_2 %>%  
  pull_workflow_fit() %$%  
  fit %>%  
  rpart.plot()
```

Tree variability

So... which should we believe?



Tree variability

Let's take several **subsamples** of the data, and make trees from each.

Then, to classify a new observation, we run it through *all* the trees and let them vote!

(It's a bit like a KNN for trees!)

This is called **bagging**

Bagging

```
library(bagette)

bag_tree_spec <- bag_tree() %>%
  set_engine("rpart", times = 5) %>%
  set_mode("classification")
```

Bagging

```
bag_tree_wflow <- workflow() %>%  
  add_recipe(cann_recipe) %>%  
  add_model(bag_tree_spec)  
  
bag_tree_fit <- bag_tree_wflow %>%  
  fit(cann)  
  
## (this code may take a while!)
```


Bagging

What variables were most important to the trees?

```
bag_tree_fit %>% pull_workflow_fit()
```

Random Forests

Random Forests

What if some important variables are being masked by *more important* variables?

Remember, we have **65** predictors - yikes!

Let's give some of the other predictors a chance to shine.

Random Forests

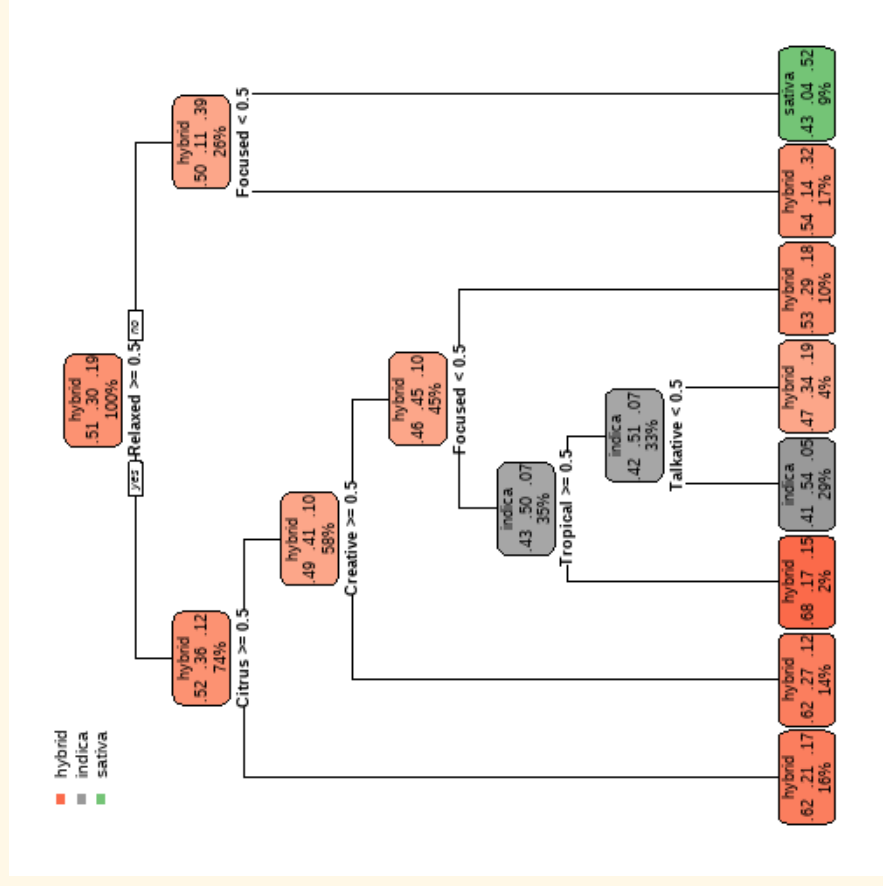
Randomly choose a few of the predictors to include in the data:

```
cann_reduced <- cann %>%  
  select(1,2, sample(5:65, 30))
```

cann_reduced

```
## # A tibble: 2,305 x 32  
##   Strain      Type Tobacco Dry Honey Ammonia Giggly Apple Relaxed Flowery Grape  
##   <chr>    <fct>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 100-0g    hybr~      0      0      0      0      0      1      0      0  
## 2 98-Whit~  hybr~      0      0      0      0      0      1      1      0  
## 3 1024      sati~      0      0      0      0      0      1      0      0  
## 4 13-Dawgs  hybr~      0      0      0      0      0      1      0      1  
## 5 24K-Gold  hybr~      0      0      0      0      0      1      0      0  
## 6 3-Bears~  indi~      0      0      0      0      0      0      0      0  
## 7 3-Kings   hybr~      0      0      0      0      0      1      0      0  
## 8 303-0g    indi~      0      0      0      1      0      1      0      0  
## 9 3D-Cbd    sati~      0      0      0      0      0      1      1      0  
## 10 3X-Crazy  indi~      0      0      0      0      0      1      0      1  
## # ... with 2,295 more rows, and 21 more variables: Tea <dbl>, Blueberry <dbl>,  
## # Creative <dbl>, Menthol <dbl>, Tropical <dbl>, Woody <dbl>, Earthy <dbl>,  
## # Skunk <dbl>, Citrus <dbl>, Mouth <dbl>, Berry <dbl>, Pungent <dbl>,  
## # Chestnut <dbl>, Blue <dbl>, Happy <dbl>, Vanilla <dbl>, Focused <dbl>,
```

Random Forests



Random Forests

After making many random reduced trees, we then **bag** the results to end up with a **random forest**.

The advantage of this is that more unique variables are involved in the process.

This way, we don't accidentally overfit to a variable that *happens* to be extremely relevant to our particular dataset.

Model spec: `rand_forest()`



Your turn

Open the activity file

Find the best *bagged* model for the cannabis data

Find the best *random forest* model for the cannabis
data

Report some metrics on your models